

11 Databricks Workflow & Automation

Databricks Workflow & Automation: Jobs, Scheduling, Monitoring, Alerts & Workflow Automation

Basic Level Questions (1-10)

1. What are Databricks Jobs and how do they differ from interactive notebooks? *Focus on: Production workloads, automation, resource management, reliability*
2. What are the different types of tasks you can create in a Databricks Job? *Focus on: Notebook tasks, JAR tasks, Python wheel tasks, SQL tasks, pipeline tasks*
3. How do you schedule a Databricks Job to run at specific intervals? *Focus on: Cron expressions, time zones, recurring schedules, trigger types*
4. What is the Databricks Workflows feature and what problems does it solve? *Focus on: Multi-task orchestration, dependency management, parallel execution, error handling*
5. How do you monitor the execution status of Databricks Jobs? *Focus on: Job runs page, execution history, logs, status indicators*
6. What are the different cluster types available for running Databricks Jobs? *Focus on: Job clusters, all-purpose clusters, pools, auto-scaling*
7. How do you pass parameters to a Databricks Job at runtime? *Focus on: Job parameters, notebook widgets, environment variables, command-line arguments*
8. What is the difference between a job cluster and an all-purpose cluster? *Focus on: Cost optimization, isolation, performance, lifecycle management*
9. How do you set up basic email notifications for job failures in Databricks? *Focus on: Notification settings, email alerts, webhook integration, escalation*
10. What are some common reasons why Databricks Jobs might fail? *Focus on: Resource issues, code errors, data problems, cluster startup failures*

Intermediate Level Questions (11-20)

11. How would you design a multi-task workflow with dependencies between different data processing stages? *Focus on: Task dependencies, conditional execution, error propagation, parallel vs sequential execution*
12. Explain how to implement retry logic and error handling in Databricks Workflows. *Focus on: Retry policies, exponential backoff, circuit breakers, dead letter queues*
13. How do you implement dynamic parameter passing between tasks in a workflow? *Focus on: Task values, parameter inheritance, dynamic configuration, state management*
14. What are the best practices for managing secrets and credentials in Databricks Jobs? *Focus on: Secret scopes, Azure Key Vault integration, environment variables, security*
15. How would you implement a data quality checkpoint system within a Databricks Workflow? *Focus on: Validation tasks, conditional branching, data quality metrics, failure handling*
16. Explain how to set up comprehensive monitoring for production Databricks Workflows. *Focus on: Metrics collection, custom logging, observability, dashboard creation*
17. How do you handle resource optimization for workflows with varying compute requirements? *Focus on: Dynamic cluster sizing, job pools, spot instances, cost optimization*
18. What strategies would you use to implement blue-green deployments for Databricks Jobs? *Focus on: Version management, deployment strategies, rollback procedures, testing*
19. How do you implement workflow orchestration for real-time and batch processing combined? *Focus on: Hybrid architectures, streaming integration, scheduling coordination, data consistency*

20. Explain how to set up custom alerting based on job performance metrics and SLA requirements. *Focus on: Custom metrics, threshold-based alerts, escalation procedures, integration with external systems*

Advanced/Difficult Level Questions (21-30)

21. Design a comprehensive CI/CD pipeline for Databricks Workflows that includes testing, deployment, and rollback capabilities. *Focus on: Infrastructure as code, automated testing, deployment automation, version control integration*

22. How would you implement a sophisticated workflow orchestration system that handles complex data lineage and dependency management across 100+ jobs? *Focus on: Dependency resolution, DAG optimization, circular dependency detection, scalability*

23. Design a multi-tenant workflow automation system in Databricks with proper isolation and resource management. *Focus on: Workspace isolation, resource quotas, access control, cost allocation, governance*

24. How would you implement intelligent job scheduling that adapts to data availability and system load? *Focus on: Event-driven scheduling, load balancing, resource optimization, predictive scheduling*

25. Explain how you would build a workflow recovery system that can automatically handle partial failures and resume from checkpoints. *Focus on: State management, checkpoint strategies, idempotency, partial execution recovery*

26. Design a comprehensive monitoring and observability solution for a complex Databricks environment with hundreds of workflows. *Focus on: Metrics aggregation, distributed tracing, log correlation, anomaly detection, predictive monitoring*

27. How would you implement a workflow optimization system that automatically tunes cluster configurations based on historical performance data? *Focus on: Machine learning for optimization, performance profiling, automated tuning, cost-performance balance*

28. Design a disaster recovery strategy for critical Databricks Workflows with RPO/RTO requirements. *Focus on: Cross-region replication, backup strategies, failover procedures, data consistency*

29. How would you implement a governance framework for workflow automation that ensures compliance and auditability? *Focus on: Policy enforcement, audit trails, compliance reporting, access controls, data governance*

30. Design a self-healing workflow system that can automatically detect, diagnose, and remediate common failure patterns. *Focus on: Failure pattern recognition, automated remediation, machine learning for anomaly detection, self-optimization*

Practical Workflow Scenarios

Real-World Implementation Challenges

Scenario A: You need to orchestrate a complex ETL pipeline that processes data from 10 different sources with varying SLAs and dependencies. How would you design this in Databricks Workflows?

Scenario B: A critical workflow is experiencing intermittent failures with no clear pattern. Walk through your troubleshooting and resolution approach.

Scenario C: You need to implement a workflow that processes both streaming and batch data with strict consistency requirements. How would you architect this?

Scenario D: Design a workflow system that can handle sudden spikes in data volume while maintaining cost efficiency.

Advanced Configuration & Integration

Enterprise Integration Patterns

External System Integration:

- REST API integration patterns
- Message queue integration (Kafka, Service Bus)
- Database connectivity and transaction management
- File system and cloud storage integration

Advanced Scheduling:

- Event-driven triggering mechanisms
- Cross-workflow dependencies
- Time-based and data-based triggers
- Complex conditional execution logic

Security & Governance:

- Role-based access control implementation
- Audit logging and compliance reporting
- Secret management and rotation
- Network security considerations

Performance & Optimization

Workflow Performance Tuning

Resource Optimization:

- Cluster sizing strategies for different workload types
- Spot instance utilization in production workflows
- Memory and CPU optimization techniques
- Storage optimization for intermediate data

Execution Optimization:

- Parallel task execution strategies
- Bottleneck identification and resolution
- Cache utilization in multi-task workflows
- Data locality optimization

Cost Management:

- Cost monitoring and attribution
- Resource usage optimization
- Automatic scaling policies
- Budget controls and alerts